

EXECUTIVE KNOWLEDGE EDITION

What Is Amsterdam Already Selling? Mapping the OTA and Competitive Tourism Product Landscape

A clean offline edition from the Ahmed Mahmoud tourism operations portfolio.

ahmedqualityops.com

What Is Amsterdam Already Selling?

Mapping the OTA and Competitive Tourism Product Landscape

Part 02 of a 12-part Amsterdam tourism product-management case study: a provenance-controlled pilot mapping public OTA listings, product families, direct competitors, substitutes, features and price-comparison limits.

What you will take away

Part 02 of a 12-part Amsterdam tourism product-management case study: a provenance-controlled pilot mapping public OTA listings, product families, direct competitors, substitutes, features and price-comparison limits.

- Explore The opening decision: map the shelf before naming the gap
- Explore The decision we need to make
- Explore Current hypotheses carried forward

A marketplace search result is not a demand report. It is a public shelf: useful for seeing what is offered, how products are framed and where offers overlap—but incapable of revealing bookings, conversion, commission or unmet need.

The opening decision: map the shelf before naming the gap

Amsterdam already sells canal cruises, walking tours, bike tours, food experiences, museum combinations, city passes, day excursions, workshops, nightlife and self-guided products through several channels. The product-management problem is not finding examples. It is deciding which offers compete for the same visitor job, which are substitutes, which listings describe the same underlying product, and which public fields can be compared without manufacturing certainty.

A search page makes the market look countable. Yet one operator may appear on multiple OTAs, one product may have several language or inclusion options, and a bundle may combine services that also sell separately. Counting every listing as a distinct competitor inflates supply. Treating a high review count as booking volume invents private performance. Treating a promoted rank as organic preference ignores the platform's own commercial logic.

This article therefore makes a narrower decision: is the competitive taxonomy and collection method strong enough to authorize the full 60–100 listing sample defined in Article 01? It does not decide where the opportunity is. That belongs to Article 03 after supply is combined with demand, customer, feasibility and destination evidence.

The decision we need to make

The output must answer four management questions:

- Which product families compete directly for a guided or self-guided Amsterdam discovery job?
- Which substitutes compete for the same time, attention or budget without delivering the same format?
- Which publicly observable features distinguish products?
- Can the sample be expanded compliantly and consistently enough for gap analysis?

The decision is not “which tour should we copy?” Competitive research should expose the grammar of the market—format, promise, duration, group model, language, inclusion, meeting logic and cancellation—not turn popular listings into a design brief.

This continues the evidence-first method from Article 01’s Amsterdam research plan. The broader product-development controls are explained in the Tourism Product Development AI Toolkit.

Current hypotheses carried forward

H01–H06 remain in the single project register. Competitive observations can add evidence, but they cannot validate customer value, willingness to pay, route acceptability, qualified guide capacity, margin or stakeholder permission.

What counts as a competitor?

A direct competitor solves a similar job through a similar service model. If the working job is “help me understand Amsterdam in a limited first-visit window,” a guided highlights walk, orientation bike tour or narrated small-group cruise may be direct even though the transport mode differs. A specialist food experience becomes direct when the proposed value is local cultural interpretation rather than transport alone.

A substitute competes for the same time or budget through a different mechanism. A city card, museum-and-cruise bundle, comedy show, museum visit, hop-on service, self-guided PDF or regional day excursion may replace the time available for a guided experience. A complement, such as a museum ticket used after a walk, can also become a substitute when the visitor has only three hours.

Directness is therefore not a permanent label. It depends on the visitor job, trip stage, duration, budget and operating area. The downloadable Direct vs Substitute Competitor Map preserves that reasoning instead of classifying products only by category name.

Research design: pilot first, census second

The target population is publicly bookable Amsterdam-area visitor experiences that could compete for discovery time and budget. The Article 01 target remains 60–100 listings. This article completed a purposive 24-record pilot across GetYourGuide, Viator, Klook, Tiqets, Expedia Experiences, I amsterdam and a direct operator page. The pilot was designed to test fields, category boundaries, duplication and access—not to estimate market share.

Tripadvisor could not be reviewed through the research tool because access was blocked. The correct response is not to circumvent the control. It is to mark the source unavailable, consider permitted manual browser review or licensed access, and test whether conclusions remain stable without it.

Publicly observable versus private or unknown

GetYourGuide publicly explains that activity ranking can consider popularity, ratings, performance, availability, diversity and revenue, including supplier commissions. Viator pages may also disclose that featured experiences can earn the platform higher commission. Ranking is therefore an outcome of several commercial and operational signals—not a clean demand variable.

Data model and provenance

Every row uses the Article 01 pipeline: RAW → STAGING → CLEAN → ANALYTICAL → DECISION. The raw title, displayed price and page wording remain untouched. Clean fields standardize duration in minutes, category, competitor type, group model, language notation and missingness. Analytical fields create product-family IDs and feature flags. The decision layer uses those fields only after a reviewer approves duplicate links and transformations.

The Competitive Product Dataset records source, URL, collection date, original price, cleaned price, price basis, duration, group model, language, rating scale, review count, itinerary, inclusions, meeting point, cancellation, positioning, differentiators, duplicate note and confidence. A missing field is not zero. “Not displayed” is different from “not offered.”

Cross-platform duplication is an analytical issue, not a reason to delete evidence. The same underlying service can have different titles, option bundles, currencies, review systems and policies. Each listing stays as an observation; a product-family key connects probable equivalents. Uncertain matches remain unresolved until operator, meeting point, duration and inclusions align.

What the 24-record pilot reveals

Observed Data — Medium confidence: the pilot shows a broad but repetitive public shelf. Canal-only products occupy several positions, from high-capacity audio cruises to live-hosted small-group and private boats. Guided walking and cycling offers often promise highlights plus “hidden gems” or local perspective. Food products combine tastings with neighbourhood interpretation, walking, cycling or boating. Bundles connect museums, cruises and transport. Self-guided products remove fixed pace and guide capacity.

Derived observation — Low confidence: differentiation appears to be assembled from a limited set of levers: mode, group intimacy, live versus recorded interpretation, language, inclusion, time of day, neighbourhood, food and drink, attraction access, flexibility and local-host positioning. That is a useful feature grammar, but the purposive pilot cannot estimate the proportion of the full market using each lever.

Duration creates recognizable shelves: roughly one-hour cruises and tastings; two-to-three-hour walks and bike tours; three-to-four-hour food or hybrid experiences; and full-day combinations or regional excursions. These bands compete for different itinerary slots. Comparing a 60-minute cruise with an 8.5-hour walk-bike-cruise product as if price alone determined value would be meaningless.

Price also requires disciplined bases. The pilot contains per-person prices, a per-group private price, promotional crossed-out prices, multiple currencies, bundles and pages where price depends on date or option. No market average is reported. Before Article 03, the full sample needs one observation date, common participant profile, common service date, retained original currency and a documented foreign-exchange date if normalization is necessary.

Category and subcategory architecture

The pilot taxonomy separates mode from experience promise. “Boat” describes delivery; “historic orientation,” “food and drink,” “evening atmosphere” or “private local interpretation” describes the job and positioning. Combining them produces useful subcategories without treating marketing phrases as unique markets.

- Canal: classic audio, live-hosted, small-group, private, food/drink, evening and attraction bundle.
- Walking: highlights, history, neighbourhood, Anne Frank/history, food and private orientation.
- Cycling: city highlights, specialist-language, neighbourhood and countryside dispersal.
- Hybrid: walk + cruise, bike + cruise, walk + bike + cruise, food + cruise.
- Independent: self-guided food, city card, attraction bundle and flexible discovery.
- Substitute leisure: museum, performance, workshop and regional day excursion.

This structure prevents two errors. First, it does not label every canal product identical. Second, it does not call every different category non-competitive. A premium small-group boat and a private historical walk may compete strongly for a couple seeking a personal two-hour orientation.

Feature matrix: compare promises, not adjectives

Words such as authentic, unique, local and unforgettable are too elastic for reliable comparison. The Competitor Feature Matrix uses inspectable features: live guide, audio or digital interpretation, maximum group where stated, private option, food/drink, admission, transport mode, language, fixed meeting point, weather exposure, accessibility disclosure and cancellation.

A feature appearing often is not automatically important. It may be easy to copy, required by the mode, or heavily represented because one platform category dominated the pilot. A feature appearing rarely is not automatically a gap. It may lack demand, be difficult to operate, conflict with destination priorities or simply use different wording.

Analytics used in this article

The pilot supports taxonomy and collection decisions. It does not yet support statistical price conclusions. Median, quartiles and box plots become appropriate only after incompatible price bases are separated.

How we should visualize this decision

- Category distribution: horizontal bars for validated product families, not raw listings. It asks where observed supply concentrates; platform sampling can distort the answer.
- Feature matrix: product families by observable feature. It shows repeated combinations; absence may mean undisclosed rather than unavailable.
- Direct–substitute map: experience similarity against competition for visitor time. It clarifies competitive logic; positions change with the visitor job.

- Price–duration plot: comparable per-person price against duration, with group model encoded. It can expose shelves; mixed currencies, dates and bundles must be excluded.

Copy-ready AI prompts for controlled competitive analysis

Normalize the competitive dataset

Use after manual collection.

Purpose: standardize fields without overwriting observations. Required input: raw rows and dictionary. Expected output: clean table plus transformation log. Human verification: analyst checks every changed field.

Competitive data normalization

Propose duplicate product families

Use before counting competitors.

Purpose: distinguish marketplace listings from underlying products. Required input: normalized rows. Expected output: candidate match groups with reasons. Human verification: researcher confirms each link.

Duplicate-linkage review

Build an auditable positioning map

Use only after taxonomy review.

Purpose: compare observable propositions without declaring a market gap. Required input: approved product families and coding rules. Expected output: positions, evidence and uncertainty. Human verification: product lead reviews axes and implications.

Positioning-map analysis

Competitive artifact pack

- Competitive Product Dataset — 24 pilot observations with provenance; full sample awaiting data.
- Competitor Feature Matrix — observable feature definitions and pilot coding.
- Direct vs Substitute Competitor Map — visitor-job-based competition logic.
- Cumulative Hypothesis Register and Assumption Register — Article 01 controls preserved and reviewed.

What we know now

- Public Amsterdam supply spans several modes but reuses a compact set of positioning and service-design levers.
- Listings are not the same as underlying products; cross-platform and option duplication must be controlled.

- Canal, walking, cycling, food, hybrid, independent and substitute offers compete differently depending on the visitor job.
- Public pages support structured supply analysis but cannot reveal demand, conversion, margin or destination permission.

What we still don't know

We do not know reliable category market shares, realized prices, bookings, conversion, willingness to pay, repeat purchase, operational quality, margins or whether any low-density position represents an attractive gap. The full 60–100 listing sample, common price protocol and duplicate review remain open.

Assumptions updated

A01 remains OPEN with additional pilot evidence: useful public evidence can be collected manually from accessible pages, but platform access is uneven and automation has not been authorized. A02–A05 remain OPEN because competitive supply does not validate recruitment, stakeholder participation, route acceptability or destination contribution.

Hypotheses updated

H01–H06 remain UNTESTED. The pilot adds supply-side observations to H01–H04 and a price-collection method to H05, but none meets its validation test. H06 receives no supporting evidence from commercial listings.

Product decision

AUTHORIZE THE FULL COMPETITIVE COLLECTION — DO NOT DECLARE A MARKET GAP. The pilot taxonomy, provenance controls and manual method are workable. Expand to the pre-agreed 60–100 listings, fix the price-observation protocol, resolve probable duplicates and document inaccessible sources before opportunity analysis.

Decision gate status

COLLECTION METHOD APPROVED WITH CONDITIONS. Article 03's PROCEED / INVESTIGATE / STOP gate is not yet open. Competitive intensity alone cannot authorize a product.

Next in the series

Article 03 — Where Is the Real Opportunity? Turning Demand, Supply and Destination Evidence Into a Market Gap. It has not been generated.

Sources and further reading

Primary standards, institutional guidance and professional references used to support this article.

- GetYourGuide — public Jordaan activity results and ranking disclosure (checked 22 August 2026).

- GetYourGuide — walking, biking and cruising product page.
- Viator — food walking tour with canal cruise.
- Klook — Amsterdam walking tour with canal cruise.
- Tqets — Van Gogh Museum and canal-cruise bundle.
- Expedia Experiences — public Amsterdam tours results.
- I amsterdam — official City Card and ticket ecosystem.
- I amsterdam — self-guided Jordaan food experience.
- Blue Boat Company — direct canal-cruise product and price page.

Editorial note: Listings, prices, ratings, review counts and policies are time-sensitive. The pilot records public observations only. It claims no access to private OTA or operator data and authorizes no automated scraping.

Continue, apply or keep the framework.

Was this article helpful?

What best describes your experience?

Keep the useful ideas coming.

Occasional, human-reviewed updates on tourism operations, product, AI, destination management and service quality.

By subscribing, you agree to receive occasional updates. You can unsubscribe at any time. Privacy details.